

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(5)

(a)

```
1 library(ISLR)
2 attach(Default)
3 set.seed(1)
4 fit <- glm(default ~ income + balance, data = Default, family = "binomial")
5 summary(fit)
6 |
```

6:1 (Top Level) ↕

Console ~/ ↶ ↷

balance, default, income, student

```
> set.seed(1)
> fit <- glm(default ~ income + balance, data = Default, family = "binomial")
> summary(fit)
```

Call:
glm(formula = default ~ income + balance, family = "binomial",
data = Default)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
Residual deviance: 1579.0 on 9997 degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8

```
> |
```

(b)

B1

```
12 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
13 |
```

13:1 (Top Level) ↕

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

B2

```
12 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
13 fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
14 summary(fit)
15
16
```

26:1 (Top Level) ↕ R Scr

Console ~/ ↗

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> summary(fit)
```

Call:
glm(formula = default ~ income + balance, family = "binomial",
data = Default, subset = train)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.1891	-0.1573	-0.0605	-0.0226	3.6623

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.136e+01	5.982e-01	-18.987	< 2e-16 ***
income	2.265e-05	6.947e-06	3.259	0.00112 **
balance	5.530e-03	3.142e-04	17.600	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1483.83 on 4999 degrees of freedom
Residual deviance: 829.05 on 4997 degrees of freedom
AIC: 835.05

Number of Fisher Scoring iterations: 8

B3

```
19 prob <- predict(fit, new= Default[-train, ], type = "response")
20 pred <- rep("No", length(prob))
21 pred[prob > 0.5] <- "Yes"
22 |
23
24
```

22:1 (Top Level) ↕

Console ~/ ↗

```
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred <- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> |
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

B4

```
19 prob <- predict(fit, new= Default[-train, ], type = "response")
20 pred <- rep("No", length(prob))
21 pred[prob > 0.5] <- "Yes"
22
23 mean(pred != Default[-train, ]$default)
24 |
25
26
27
```

24:1 (Top Level) ↕

Console ~/ ↗

```
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred <- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> mean(pred != Default[-train, ]$default)
[1] 0.028
> |
```

(c)

Depending on different predictors included in the training and validation set, the validation test error rate is varying.

```
25 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
26 fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
27 prob <- predict(fit, new= Default[-train, ], type = "response")
28 pred<- rep("No", length(prob))
29 pred[prob > 0.5] <- "Yes"
30 mean(pred != Default[-train, ]$default)
31 |
```

31:1 (Top Level) ↕ R

Console ~/ ↗

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred<- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> mean(pred != Default[-train, ]$default)
[1] 0.0252
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred<- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> mean(pred != Default[-train, ]$default)
[1] 0.0246
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred<- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> mean(pred != Default[-train, ]$default)
[1] 0.0266
> |
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(d)

Including student dummy variable is not showing any decreasing difference in the test error

```
25 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
26 fit <- glm(default ~ income + balance+student, data = Default, family = "binomial", subset = train)
27 prob <- predict(fit, new= Default[-train, ], type = "response")
28 pred<- rep("No", length(prob))
29 pred[prob > 0.5] <- "Yes"
30 mean(pred != Default[-train, ]$default)
31 |
32
33
34
```

31:1 (Top Level) R Script

Console ~/

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit <- glm(default ~ income + balance+student, data = Default, family = "binomial", subset = train)
> prob <- predict(fit, new= Default[-train, ], type = "response")
> pred<- rep("No", length(prob))
> pred[prob > 0.5] <- "Yes"
> mean(pred != Default[-train, ]$default)
[1] 0.0254
> |
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(6)

(a)

The standard error rates of the coefficients beta0, beta1 and beta2 are 0.4348, 0.000004985, 0.0002274

```
1 library(ISLR)
2 attach(Default)
3 set.seed(1)
4 fit <- glm(default ~ income + balance, data = Default, family = "binomial")
5 summary(fit)
6
```

7:1 (Top Level) ↕

Console ~/ ↕

```
> fit <- glm(default ~ income + balance, data = Default, family = "binomial")
> summary(fit)
```

Call:
glm(formula = default ~ income + balance, family = "binomial",
data = Default)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16	***
income	2.081e-05	4.985e-06	4.174	2.99e-05	***
balance	5.647e-03	2.274e-04	24.836	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
Residual deviance: 1579.0 on 9997 degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8

```
> |
```

(b)

```
32 boot.fn <- function(data, index) {
33   fit <- glm(default ~ income + balance, data = data, family = "binomial", subset = index)
34   return (coef(fit))
35 }
36
```

37:14 (Top Level) ↕ R

Console ~/ ↕

```
> boot.fn <- function(data, index) {
+   fit <- glm(default ~ income + balance, data = data, family = "binomial", subset = index)
+   return (coef(fit))
+ }
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(c)

The bootstrap standard error rates of beta0, beta1 and beta2 are 4.285621e-01, 4.950151e-06 and 2.236899e-04

```
36
37 library(boot)|
38 boot(Default, boot.fn, 1000)
39
```

37:14 (Top Level) ↕

Console ~/ ↗

```
> library("boot", lib.loc="C:/Program Files/R/R-3.5.2/library")
> boot(Default, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = Default, statistic = boot.fn, R = 1000)

Bootstrap Statistics :

	original	bias	std. error
t1*	-1.154047e+01	-2.099044e-02	4.285621e-01
t2*	2.080898e-05	1.716564e-07	4.950151e-06
t3*	5.647103e-03	6.622474e-06	2.236899e-04

```
>
```

(d)

The two methods are giving approximately same error rates.

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(9)

(a)

```
58 library(MASS)
59 attach(Boston)
60 mue<-mean(medv)
61 mue
62 |
63
64
65
66
67
68
69
70
71
```

62:1 (Top Level) ↕

Console ~/ ↗

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[workspace loaded from ~/.RData]

```
> library(MASS)
> attach(Boston)
The following object is masked _by_ '.GlobalEnv':
```

age

```
> mue<-mean(medv)
> mue
[1] 22.53281
> |
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(b)

```
64 se <- sd(medv) / sqrt(dim(Boston)[1])
65 se
66 |
67
68
69
70
71
```

66:1 (Top Level) ↕

Console ~/ ↗

```
> se <- sd(medv) / sqrt(dim(Boston)[1])
> se
[1] 0.4088611
> |
```

(c)

The bootstrap standard error rate is 0.41193 which is closer to 0.4088 that is the value obtained from b

```
85 set.seed(1)
86 boot.fn <- function(data, index) {
87   mu <- mean(data[index])
88   return (mu)
89 }
90 library(boot)
91 boot(medv, boot.fn, 1000)
```

92:1 (Top Level) ↕

Console ~/ ↗

```
> set.seed(1)
> boot.fn <- function(data, index) {
+   mu <- mean(data[index])
+   return (mu)
+ }
> library(boot)
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :

	original	bias	std. error
t1*	22.53281	0.008517589	0.4119374

```
> |
```


Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(d)

The confidence interval of `t.test()` is closer to the confidence interval given by bootstrap

```
93 t.test(medv)
94 ci<- c(22.53 - 2 * 0.4119, 22.53 + 2 * 0.4119)
95 ci
96 |
```

96:1 (Top Level) ⚙

Console ~/ ↗

```
> t.test(medv)

      One Sample t-test

data:  medv
t = 55.111, df = 505, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 21.72953 23.33608
sample estimates:
mean of x
 22.53281

> ci<- c(22.53 - 2 * 0.4119, 22.53 + 2 * 0.4119)
> ci
[1] 21.7062 23.3538
> |
```

(e)

```
97 med <- median(medv)
98 med
99 < |
```

99:1 (Top Level) ⚙

Console ~/ ↗

```
> med <- median(medv)
> med
[1] 21.2
> |
```

\

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(f) The median value is 21.2 which is equal to the value obtained in (e), with a standard error of 0.3745

```
100
101 boot.fn <- function(data, index) {
102   mu <- median(data[index])
103   return (mu)
104 }
105 boot(medv, boot.fn, 1000)
106 |
<

106:1 (Top Level) ↕

Console ~/ ↗ ↶
> boot.fn <- function(data, index) {
+   mu <- median(data[index])
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original    bias      std. error
t1*      21.2 -0.0025      0.374358
> |
```

(g)

```
100
107 quantitymue <- quantile(medv, c(0.1))
108 quantitymue
109 |
<

109:1 (Top Level) ↕

Console ~/ ↗ ↶
> quantitymue <- quantile(medv, c(0.1))
> quantitymue
10%
12.75
> |
```

Santhoshini sree Bolisetty
Alexa Summers
Gireesh kumar Muppalla
Lab #3-CS 5565

(h)

The tenth percentile is 12.75 which is again equal to the value obtained in (g), with a standard error of 0.49122

```
111 boot.fn <- function(data, index) {  
112   mu <- quantile(data[index], c(0.1))  
113   return (mu)  
114 }  
115 boot(medv, boot.fn, 1000)  
116 |
```

116:1 (Top Level) ↕

Console ~/ ↗

```
> boot.fn <- function(data, index) {  
+   mu <- quantile(data[index], c(0.1))  
+   return (mu)  
+ }  
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	12.75	0.0261	0.4912231

```
> |
```