

CHAPTER-6

8.

(a)

```
1 set.seed(1)
2 x <- rnorm(100)
3 noise <- rnorm(100)
4 |
```

(b)

```
4
5 Y <- 3 + 1*X + 4*X^2 - 1*X^3 + noise
6 |
```

(c)

Regsubsets chooses three as the optimal number of parameters



Alexa Summers
Santhoshini Sree Bolisetty
Gireesh Kumar Muppalla
CS 5565—Lab 4

(d)

The backward and forward stepwise model agrees with the best subsets model

```
> model <- train(Y ~ poly(X, 10), data = df,
+               method = 'glmStepAIC', direction = 'backward',
+               trace = 0,
+               trControl = trainControl(method = 'none', verboseIter = FALSE))
> postResample(predict(model, df), df$Y)
      RMSE Rsquared      MAE
0.9314956 0.9569843 0.7488821
> summary(model$finalModel)

Call:
NULL

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.8914  -0.5860  -0.1516   0.5892   2.1794

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    6.10265    0.09557   63.856 < 2e-16 ***
poly(X, 10)1    -7.19295    0.95569  -7.526 2.96e-11 ***
poly(X, 10)2    40.74405    0.95569  42.633 < 2e-16 ***
poly(X, 10)3   -14.70908    0.95569 -15.391 < 2e-16 ***
poly(X, 10)5     1.48019    0.95569   1.549  0.125
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.9133516)

Null deviance: 2017.132 on 99 degrees of freedom
Residual deviance: 86.768 on 95 degrees of freedom
AIC: 281.59

Number of Fisher Scoring iterations: 2
```

```
> x_poly <- poly(df$X, 10)
> colnames(x_poly) <- paste0('poly', 1:10)
> model_forw <- train(Y = Y, x = x_poly,
+                    method = 'glmStepAIC', direction = 'forward',
+                    trace = 0,
+                    trControl = trainControl(method = 'none', verboseIter = FALSE))
> postResample(predict(model_forw, data.frame(x_poly)), df$Y)
      RMSE Rsquared      MAE
0.9314956 0.9569843 0.7488821
> summary(model_forw$finalModel)

Call:
NULL

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.8914  -0.5860  -0.1516   0.5892   2.1794

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    6.10265    0.09557   63.856 < 2e-16 ***
poly2          40.74405    0.95569  42.633 < 2e-16 ***
poly3         -14.70908    0.95569 -15.391 < 2e-16 ***
poly1          -7.19295    0.95569  -7.526 2.96e-11 ***
poly5           1.48019    0.95569   1.549  0.125
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

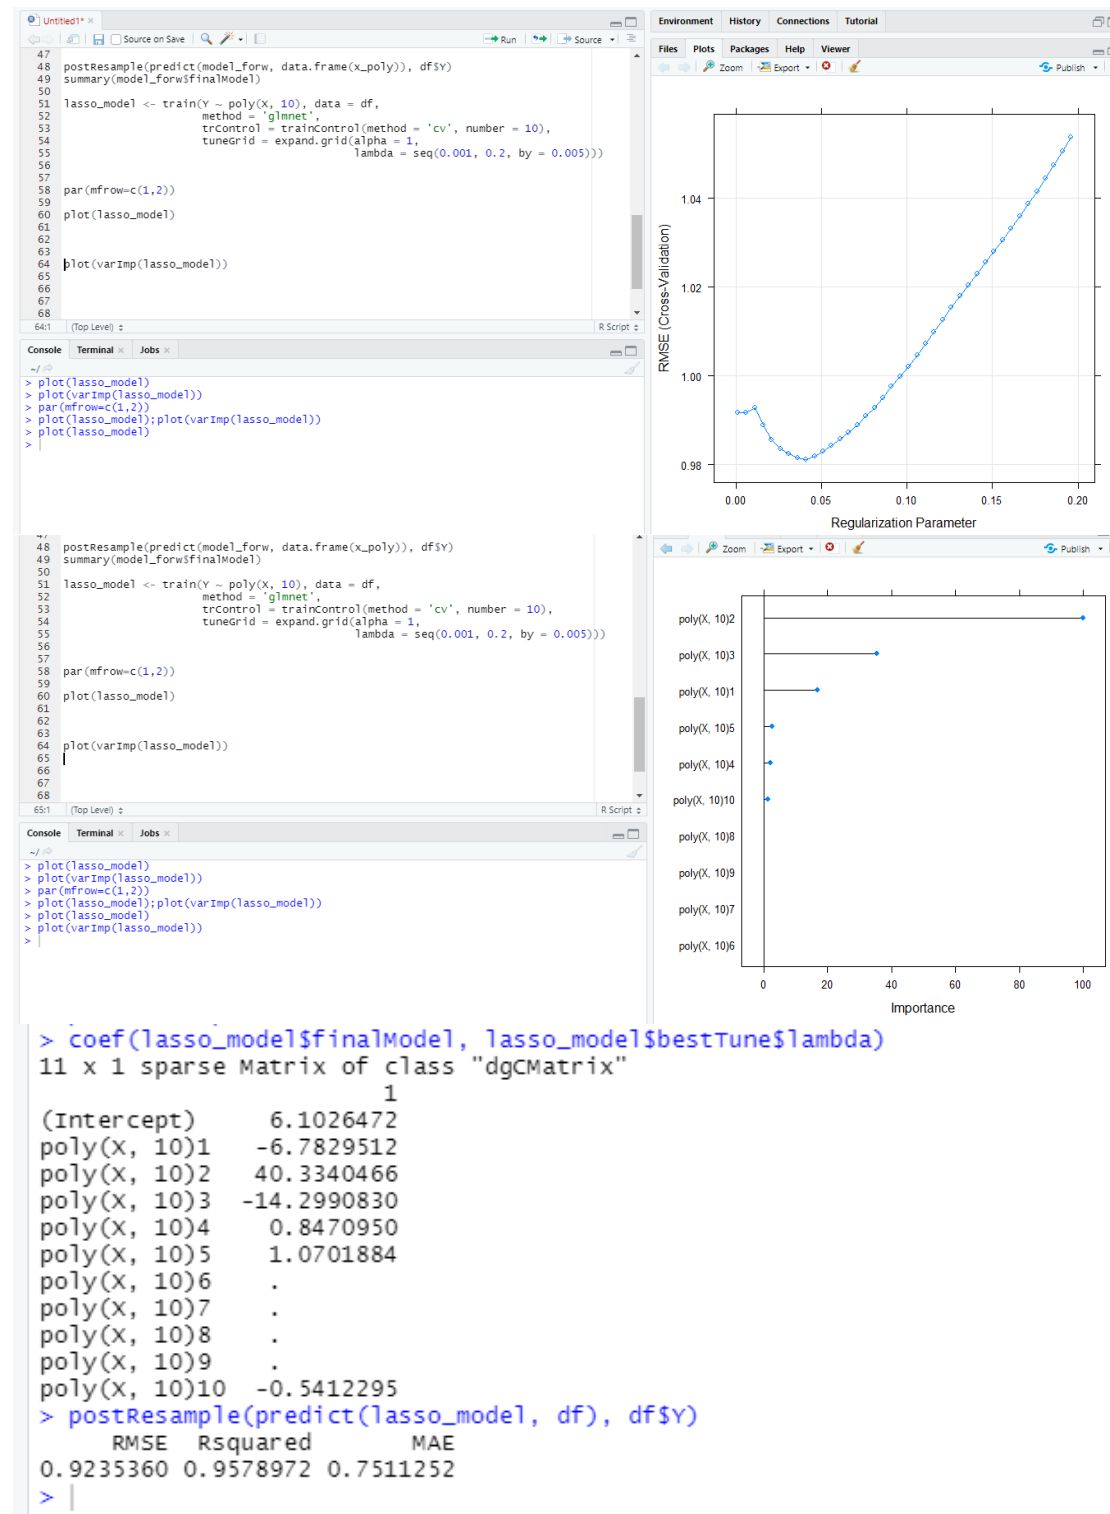
(Dispersion parameter for gaussian family taken to be 0.9133516)

Null deviance: 2017.132 on 99 degrees of freedom
Residual deviance: 86.768 on 95 degrees of freedom
AIC: 281.59

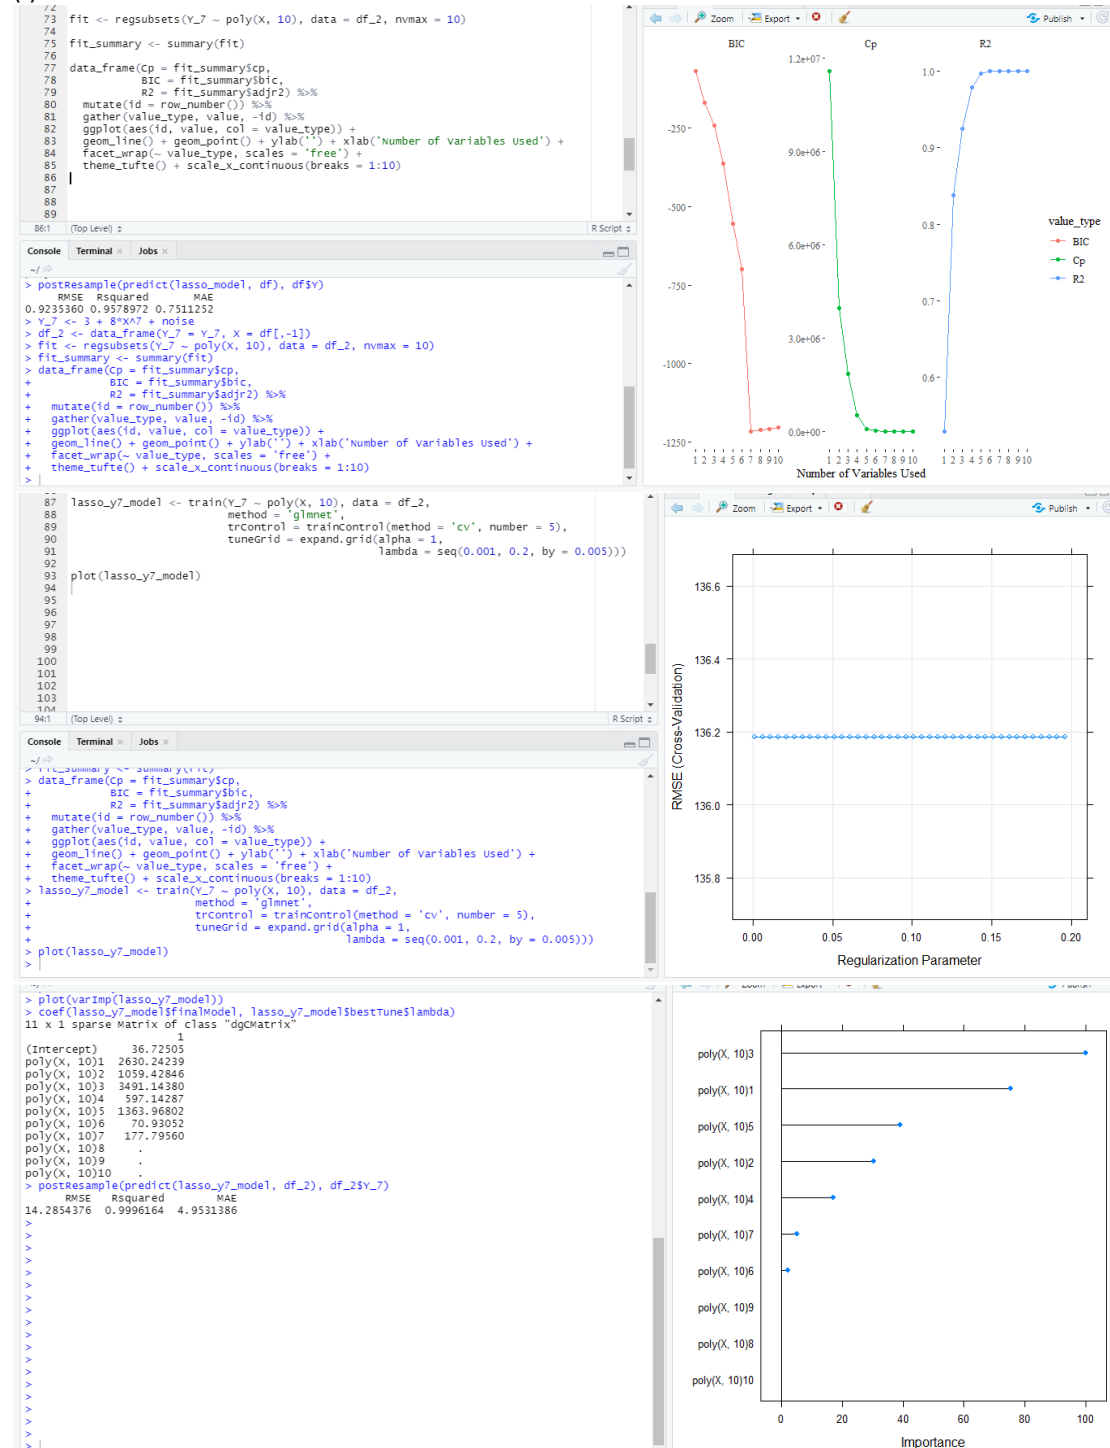
Number of Fisher Scoring iterations: 2
```

(e)

The Lasso model overestimates the number of predictors needed. This might be expected since we used only RSS to select the optimal model but not the Bayesian Inference Criterion and Adjusted R^2 as regsubsets does or the Aikake Information Criterion as the stepwise selection does.



(f)



9.

(a)

```
library(ISLR)
library(caret)
library(tidyverse)
data('College')
set.seed(1)

inTrain <- createDataPartition(College$Apps, p = 0.75, list = FALSE)

training <- College[inTrain,]
testing <- College[-inTrain,]

obj <- preProcess(training, method = c('center', 'scale'))

training <- predict(obj, training)
testing <- predict(obj, testing)

y_train <- training$Apps
y_test <- testing$Apps


apps <- dummyVars(Apps ~ ., data = training)
x_train <- predict(apps, training)
x_test <- predict(apps, testing)
```

(b)

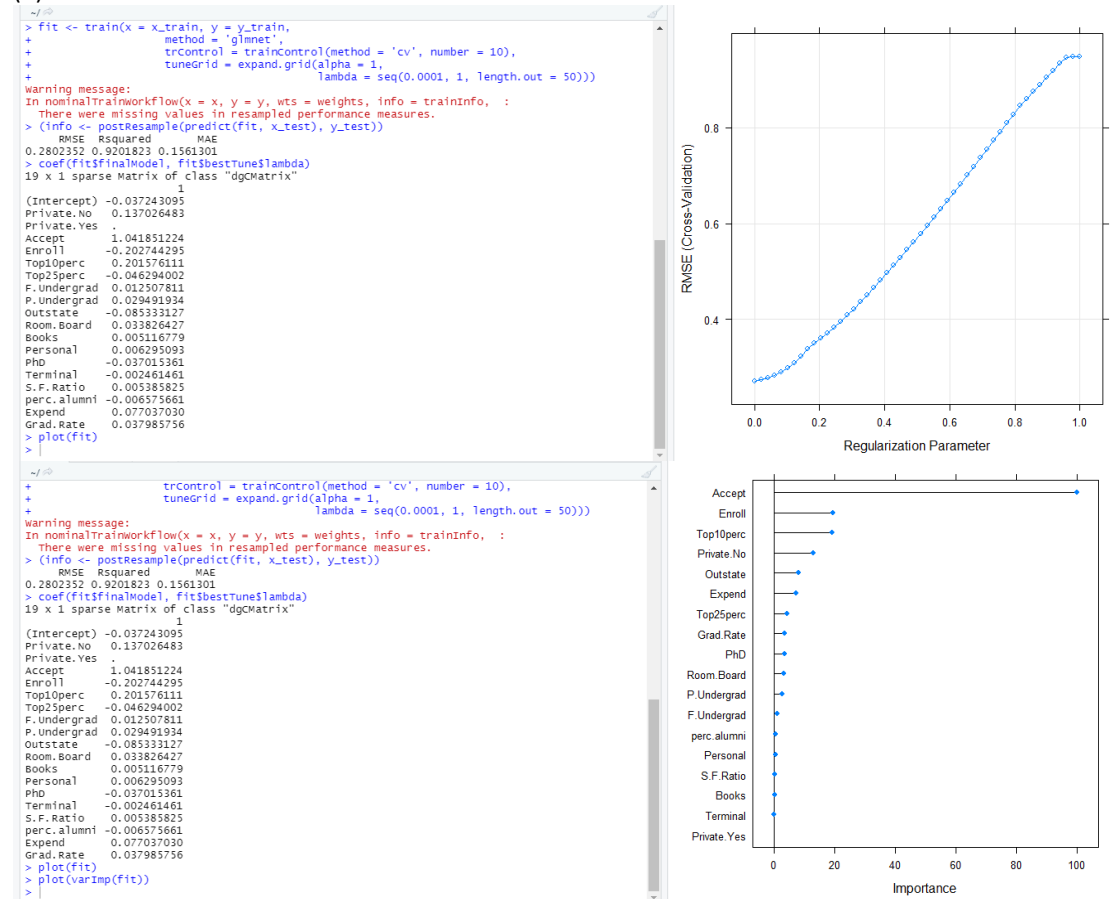
```
model <- lm(Apps ~ ., data = training)
pred <- predict(model, testing)
info <- postResample(pred, testing$Apps)
```

(c)

```
> fit <- train(x = x_train, y = y_train,
+             method = 'glmnet',
+             trControl = trainControl(method = 'cv', number = 10),
+             tuneGrid = expand.grid(alpha = 0,
+                                   lambda = seq(0, 10e2, length.out = 20)))
Warning message:
In nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
  There were missing values in resampled performance measures.
> (info <- postResample(predict(ridge_fit, x_test), y_test))
      RMSE Rsquared      MAE
0.2853247 0.9211286 0.1645806
> coef(fit$finalModel, fit$bestTune$lambda)
19 x 1 sparse Matrix of class "dgcmatrix"
      1
(Intercept) 0.034871314
Private.No 0.075423210
Private.Yes -0.076037580
Accept 0.665628733
Enroll 0.090243372
Top10perc 0.107160248
Top25perc 0.011628030
F.Undergrad 0.063308801
P.Undergrad 0.017427317
Outstate -0.028995432
Room.Board 0.048720533
Books 0.012799145
Personal -0.002894430
PhD -0.017989230
Terminal -0.010434665
S.F.Ratio 0.006920126
perc.alumni -0.031683867
Expend 0.083525070
Grad.Rate 0.058131023
> plot(ridge_fit)
> plot(varimp(ridge_fit))
>
```



(d)



(e)

```

package 'pls' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  c:\Users\Santhoshini Sree\AppData\Local\Temp\Rtmp4e4wQZ\downloaded_packages
> pcr_model <- train(x = x_train, y = y_train,
+                   method = 'pcr',
+                   trControl = trainControl(method = 'cv', number = 10),
+                   tuneGrid = expand.grid(ncomp = 1:10))
> (pcr_info <- postResample(predict(pcr_model, x_test), y_test))
      RMSE    Rsquared    MAE
0.3231292 0.8916531 0.1986075
> coef(pcr_model$finalModel)
, , 10 comps

      .outcome
Private.No    0.031985972
Private.Yes  -0.031985972
Accept       0.343576750
Enroll       0.305359773
Top10perc    0.042630417
Top25perc    0.027790893
F.Undergrad  0.273818439
P.Undergrad  -0.049487667
Outstate     0.038573119
Room.Board   0.070607615
Books        0.016433593
Personal     -0.023529455
PhD          -0.023992433
Terminal     -0.024182230
S.F.Ratio    0.003741623
perc.alumni  -0.070567887
Expend       0.090126298
Grad.Rate    0.071302714

> plot(pcr_model)
>
173:1 (Top Level) z R Script

```



```

~/.R
> pcr_model <- train(x = x_train, y = y_train,
+                   method = 'pcr',
+                   trControl = trainControl(method = 'cv', number = 10),
+                   tuneGrid = expand.grid(ncomp = 1:10))
> (pcr_info <- postResample(predict(pcr_model, x_test), y_test))
      RMSE    Rsquared    MAE
0.3231292 0.8916531 0.1986075
> coef(pcr_model$finalModel)
, , 10 comps

      .outcome
Private.No    0.031985972
Private.Yes  -0.031985972
Accept       0.343576750
Enroll       0.305359773
Top10perc    0.042630417
Top25perc    0.027790893
F.Undergrad  0.273818439
P.Undergrad  -0.049487667
Outstate     0.038573119
Room.Board   0.070607615
Books        0.016433593
Personal     -0.023529455
PhD          -0.023992433
Terminal     -0.024182230
S.F.Ratio    0.003741623
perc.alumni  -0.070567887
Expend       0.090126298
Grad.Rate    0.071302714

> plot(pcr_model)
> plot(varImp(pcr_model))
>

```



(f)


```

~/.R
> model <- train(x = x_train, y = y_train,
+               method = 'pls',
+               trControl = trainControl(method = 'cv', number = 10),
+               tuneGrid = expand.grid(ncomp = 1:10))
> (info <- postResample(predict(model, x_test), y_test))
      RMSE    Rsquared    MAE
0.2792580 0.9209302 0.1572165
> coef(model$finalModel)
, , 10 comps

      .outcome
Private.No    0.071314109
Private.Yes  -0.071314109
Accept       1.039263053
Enroll       -0.169855531
Top10perc    0.235572152
Top25perc    -0.070270182
F.Undergrad  -0.022580399
P.Undergrad  0.032522451
Outstate     -0.085919942
Room.Board   0.036344367
Books        0.004835496
Personal     0.007546695
PhD          -0.044709439
Terminal     0.002332025
S.F.Ratio    0.010868271
perc.alumni  -0.009046689
Expend       0.072963738
Grad.Rate    0.037425850

> plot(model)
>


```



(g)

```
> as_data_frame(rbind(lin_info,
+                      ridge_info,
+                      lasso_info,
+                      pcr_info,
+                      pls_info)) %>%
+   mutate(model = c('Linear', 'Ridge', 'Lasso', 'PCR', 'PLS')) %>%
+   select(model, RMSE, Rsquared)
# A tibble: 5 x 3
  model    RMSE Rsquared
  <chr>    <dbl>   <dbl>
1 Linear  0.280    0.920
2 Ridge   0.285    0.921
3 Lasso   0.291    0.914
4 PCR     0.323    0.892
5 PLS     0.279    0.921

#> sd
#> library(ggthemes)
#> residfunc <- function(fit, data) {
#>   predict(fit, data) - testing$Apps
#> }
#> data_frame(observed = testing$Apps,
#>             LM = residfunc(lin_model, testing),
#>             Ridge = residfunc(ridge_fit, x_test),
#>             Lasso = residfunc(lasso_fit, x_test),
#>             PCR = residfunc(pcr_model, x_test),
#>             PLS = residfunc(pls_model, x_test)) %>%
#>   gather(model, Residuals, -observed) %>%
#>   ggplot(aes(observed, Residuals, col = model)) +
#>   geom_hline(yintercept = 0, lty = 2) +
#>   geom_point(alpha = 0.6) +
#>   geom_smooth(method = 'loess', alpha = 0.01, col = 'lightsalmon2') +
#>   facet_wrap(~ model, ncol = 5) +
#>   theme_tufte() +
#>   theme(legend.position = 'top') +
#>   coord_flip()
#>   geom_smooth() using formula 'y ~ x'
#>   warning message:
#>   attributes are not identical across measure variables;
#>   they will be dropped
#> }
```



10.

(a)

```
> set.seed(1)
> x <- matrix(rnorm(1000 * 20), 1000, 20)
> b <- rnorm(20)
> b[3] <- 0
> b[4] <- 0
> b[9] <- 0
> b[19] <- 0
> b[10] <- 0
> eps <- rnorm(1000)
> y <- x %*% b + eps
>
```

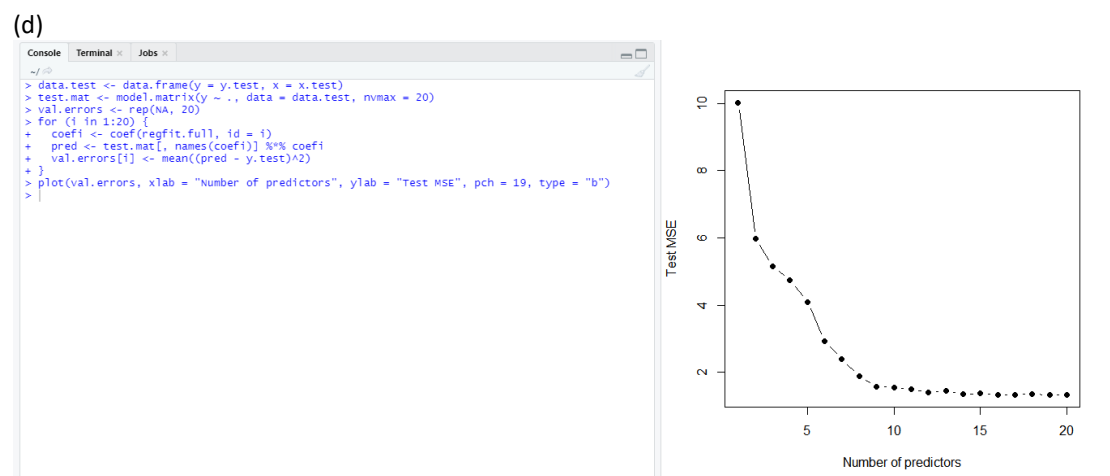
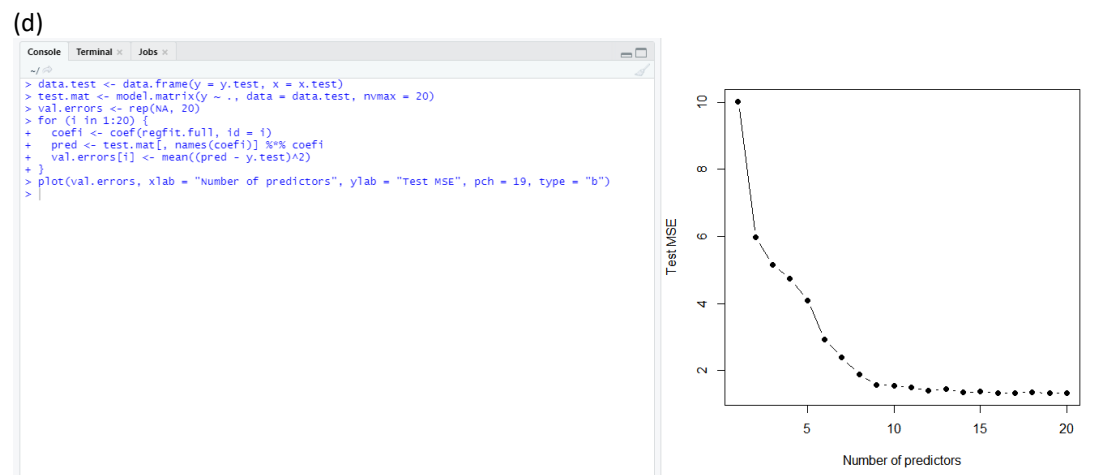
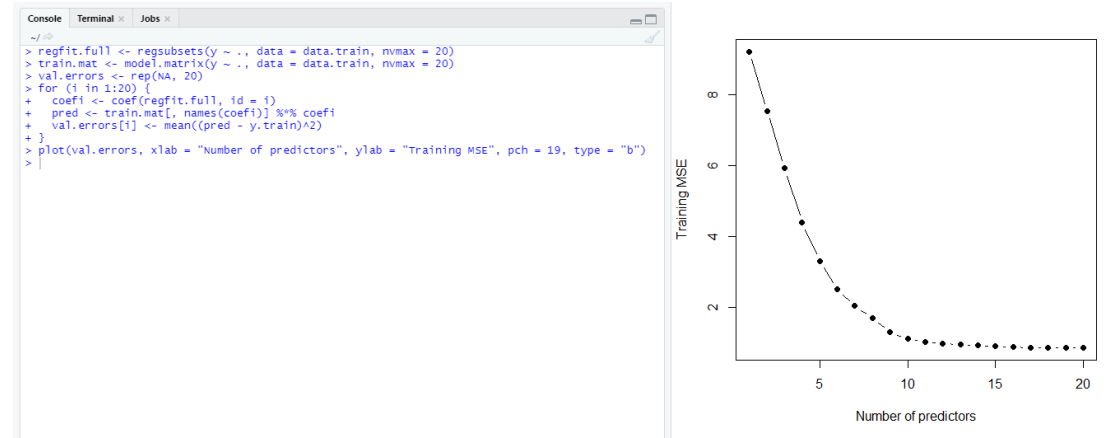
(b)

```
Console Terminal x Jobs x
~/
> train <- sample(seq(1000), 100, replace = FALSE)
> test <- -train
> x.train <- x[train, ]
> x.test <- x[test, ]
> y.train <- y[train]
> y.test <- y[test]
>
```

(c)

```
Console Terminal x Jobs x
~/
> data.train <- data.frame(y = y.train, x = x.train)
> library(leaps)
> regfit.full <- regsubsets(y ~ ., data = data.train, nvmax = 20)
> train.mat <- model.matrix(y ~ ., data = data.train, nvmax = 20)
> val.errors <- rep(NA, 20)
> for (i in 1:20) {
+   coefi <- coef(regfit.full, id = i)
+   pred <- train.mat[, names(coefi)] %*% coefi
+   val.errors[i] <- mean((pred - y.train)^2)
+ }
>
```

Alexa Summers
Santhoshini Sree Bolisetty
Gireesh Kumar Muppalla
CS 5565—Lab 4



(e)

```
> which.min(val.errors)
[1] 19
> |
```

14-variables model has the smallest test MSE.

(f)



```
> coef(regfit.full, which.min(val.errors))
(Intercept)      x.1      x.2      x.3      x.5      x.6      x.7
-0.04011778  0.13698928  0.20775097 -0.10438563  1.04174837 -0.21709681 -1.31761521
      x.8      x.9      x.10     x.11     x.12     x.13     x.14
 0.72571564  0.12387155 -0.18363204  1.01888399  0.64149490 -0.41714902 -0.70005302
      x.15     x.16     x.17     x.18     x.19     x.20
-0.76105664 -0.40810077  0.04686190  1.65688296 -0.13786948 -0.99777611
> |
```

The best model caught all zeroed out coefficients.

(g)

We may see that the model with 3 variables minimizes the error between the estimated and true coefficients. However test error is minimized by the model with 14 variables. So, a better fit of true coefficients doesn't necessarily mean a lower test MSE.

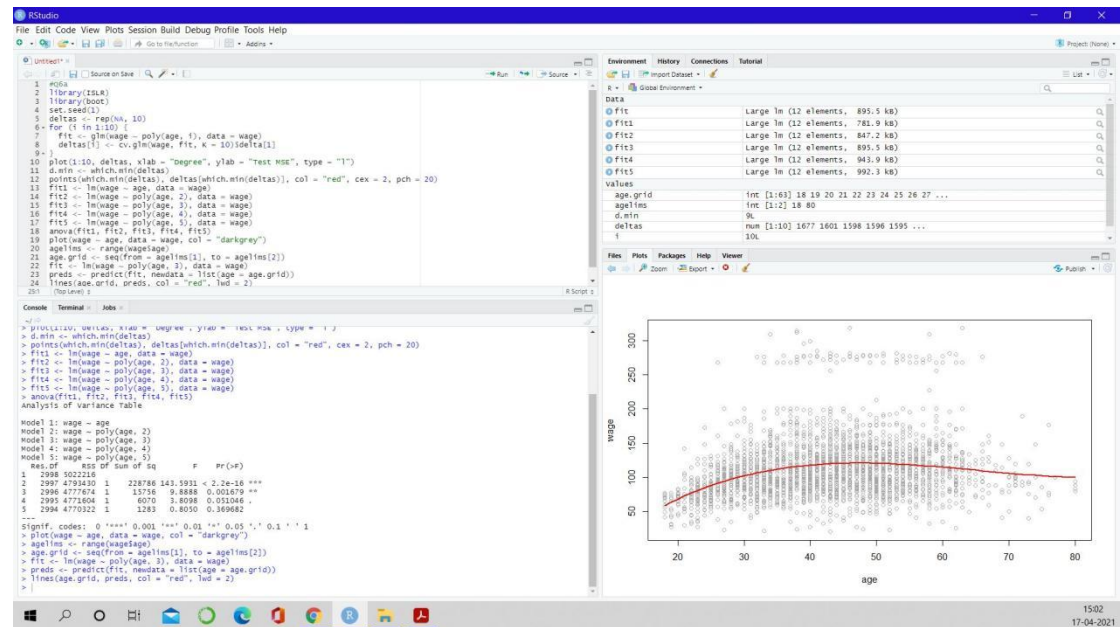
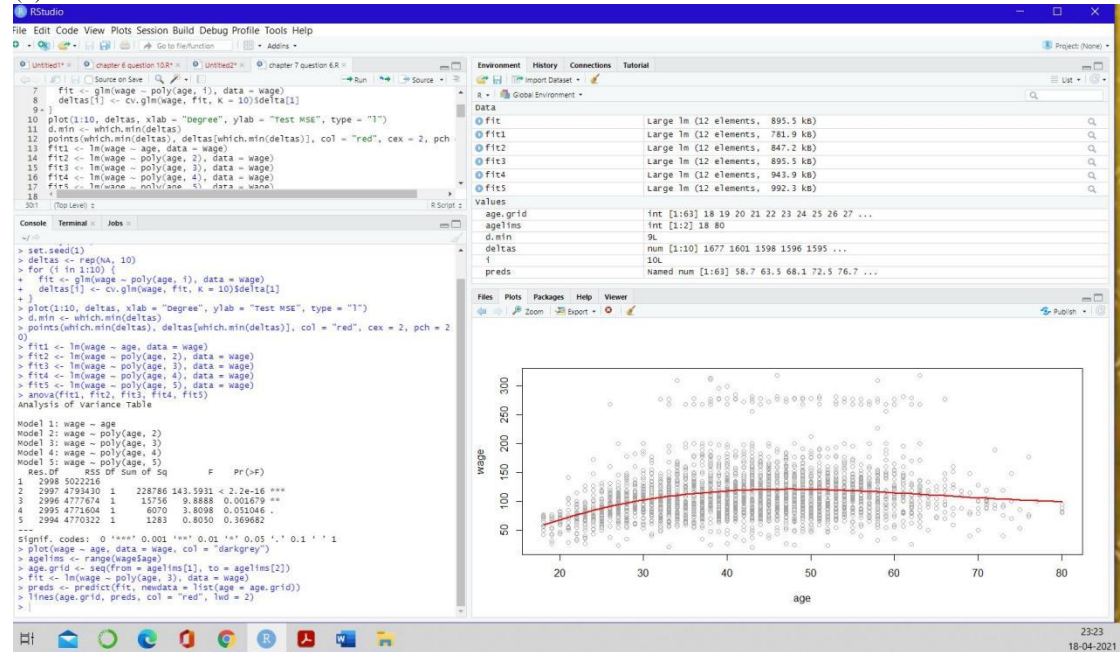
```
> data.test <- data.frame(y = y.test, x = x.test)
> test.mat <- model.matrix(y ~ ., data = data.test, nvmax = 20)
> val.errors <- rep(NA, 20)
> for (i in 1:20) {
+   coef1 <- coef(regfit.full, id = i)
+   pred <- test.mat[, names(coef1)] %*% coef1
+   val.errors[i] <- mean((pred - y.test)^2)
+ }
> plot(val.errors, xlab = "Number of predictors", ylab = "Test MSE", pch = 19, type = "b")
> which.min(val.errors)
[1] 19
> coef(regfit.full, which.min(val.errors))
(Intercept)      x.1      x.2      x.3      x.5      x.6      x.7
-0.04011778  0.13698928  0.20775097 -0.10438563  1.04174837 -0.21709681 -1.31761521
      x.8      x.9      x.10     x.11     x.12     x.13     x.14
 0.72571564  0.12387155 -0.18363204  1.01888399  0.64149490 -0.41714902 -0.70005302
      x.15     x.16     x.17     x.18     x.19     x.20
-0.76105664 -0.40810077  0.04686190  1.65688296 -0.13786948 -0.99777611
> val.errors <- rep(NA, 20)
> x.cols <- colnames(x, do.NULL = FALSE, prefix = "x.")
> for (i in 1:20) {
+   coef1 <- coef(regfit.full, id = i)
+   val.errors[i] <- sqrt(sum(b[x.cols %in% names(coef1)] - coef1[names(coef1) %in% x.cols]^2)
+   + sum(b[!(x.cols %in% names(coef1))]^2))
+ }
> plot(val.errors, xlab = "Number of predictors", ylab = "Error between estimated and true coefficients", pch = 19, type = "b")
> |
```

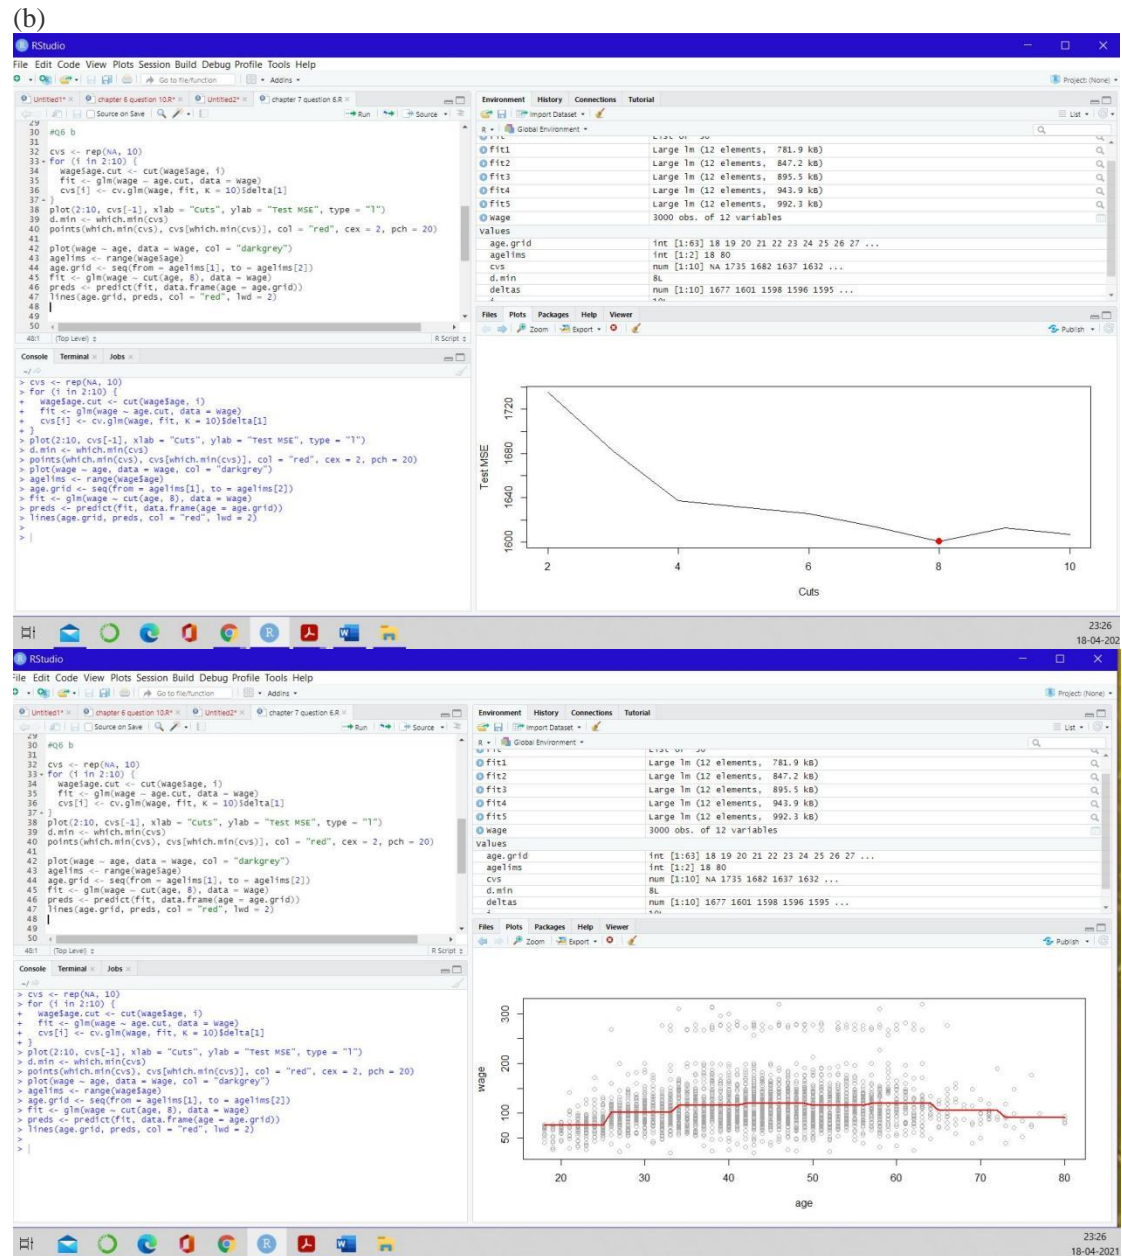


CHAPTER 7

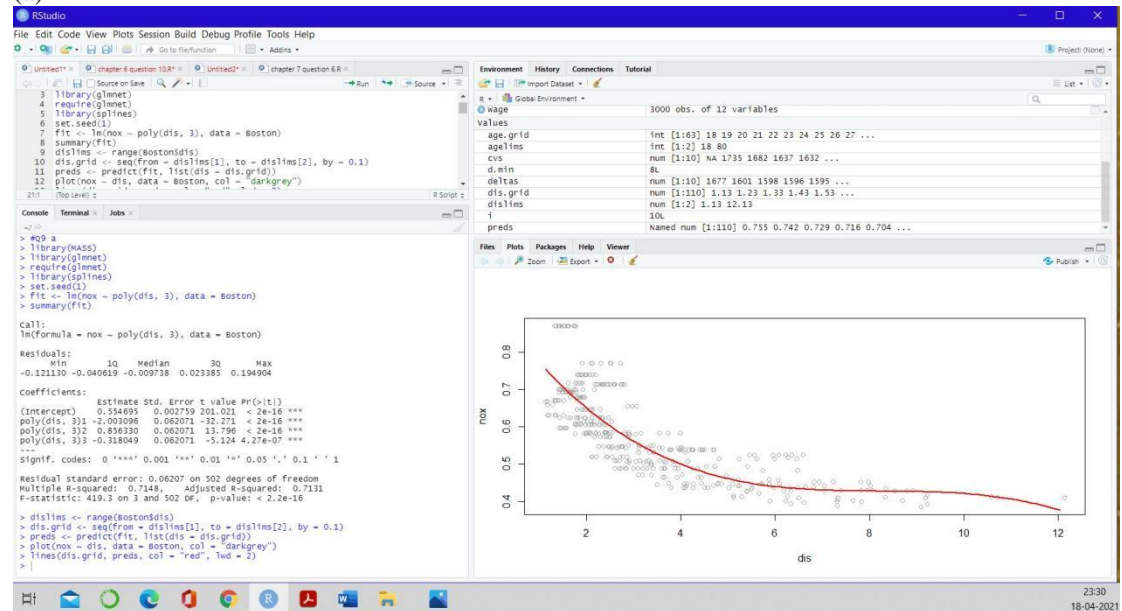
6.

(a)

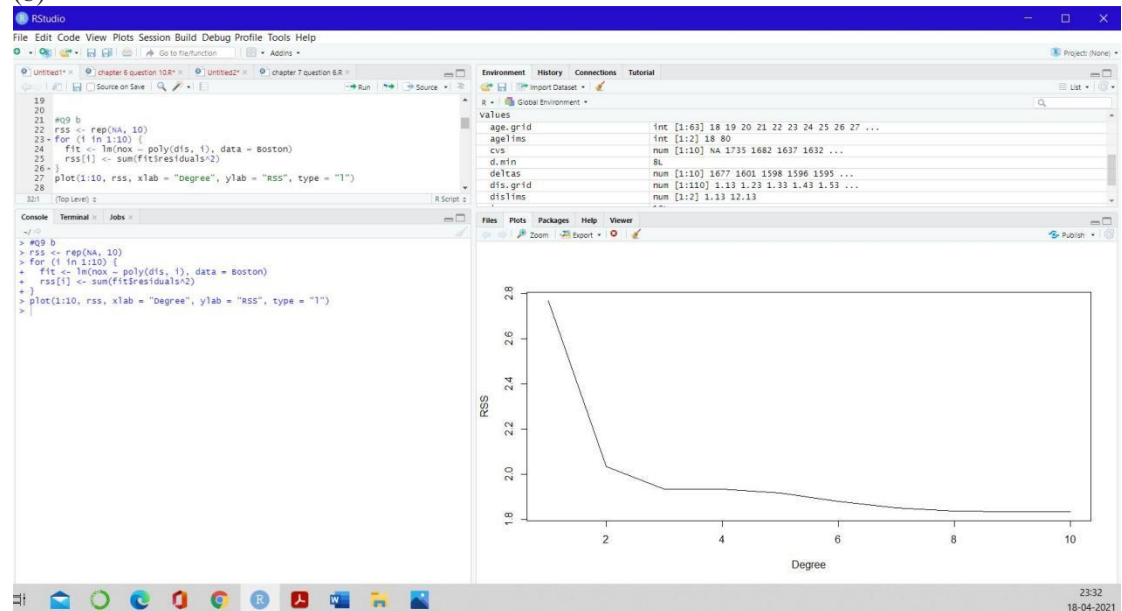




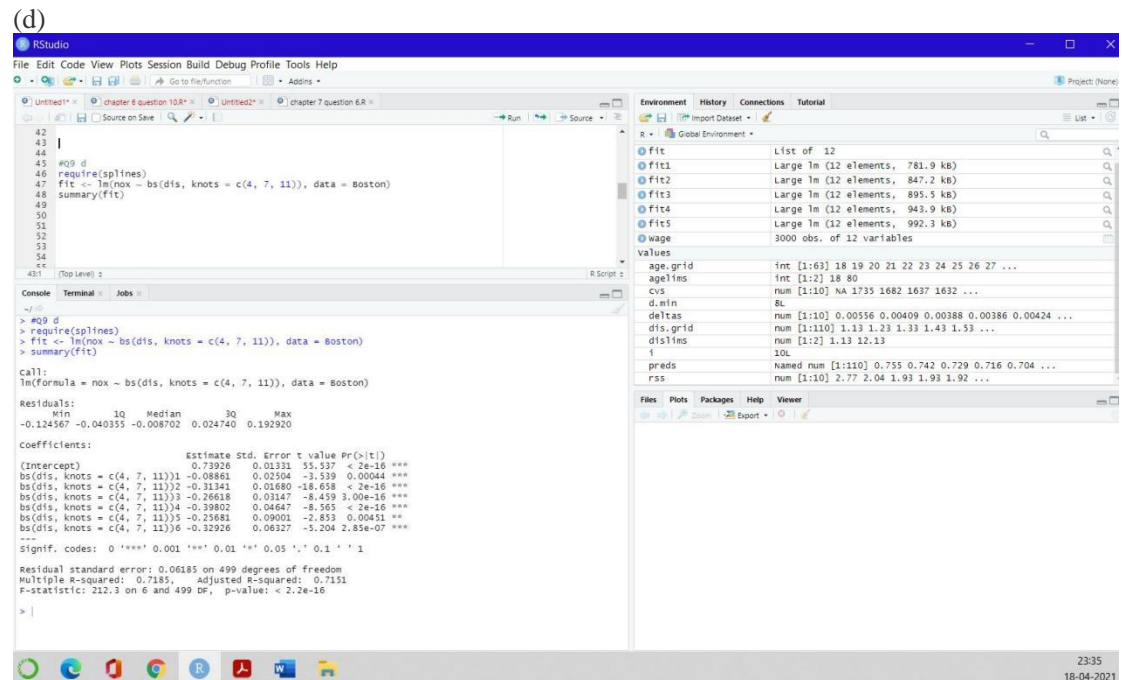
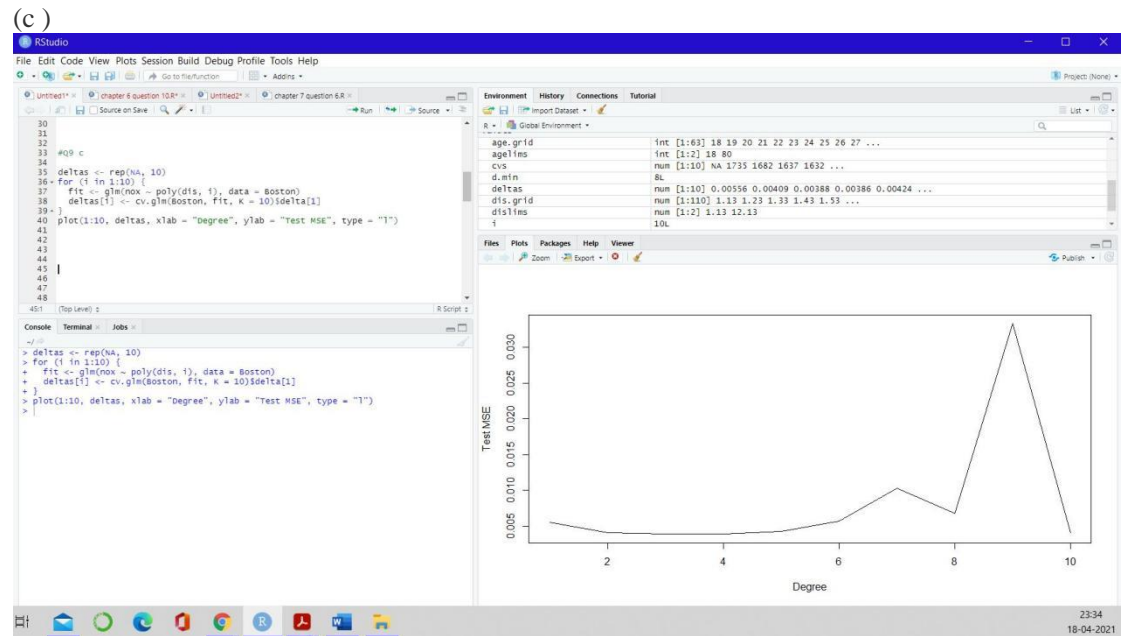
9.
 (a)



(b)

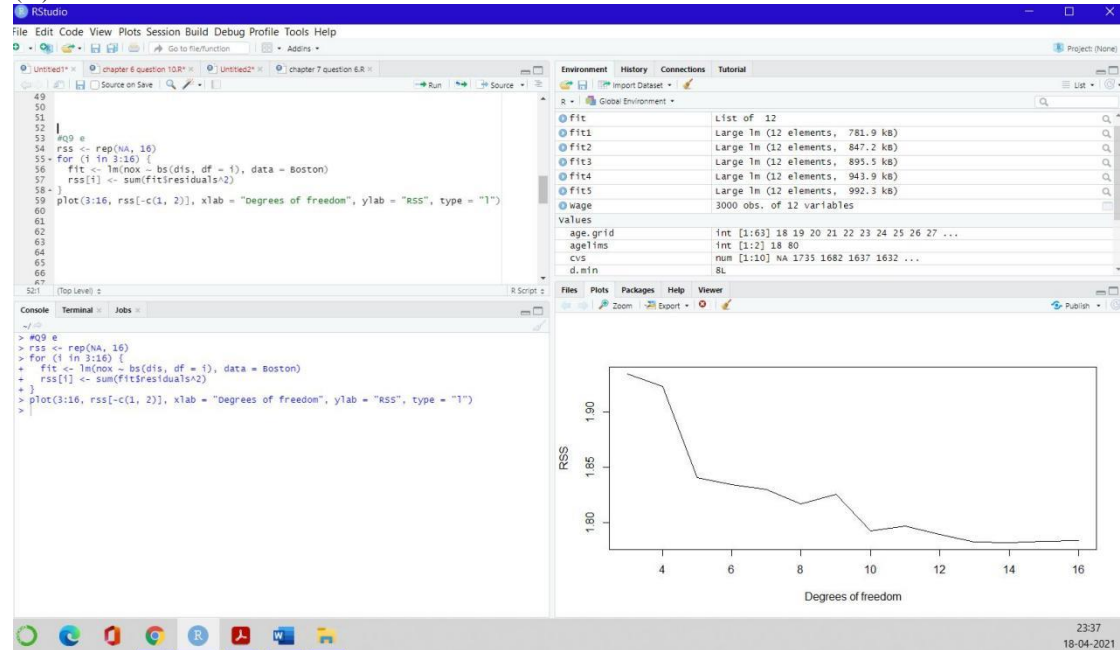


Alexa Summers
Santhoshini Sree Bolisetty
Gireesh Kumar Muppalla
CS 5565—Lab 4



Alexa Summers
Santhoshini Sree Bolisetty
Gireesh Kumar Muppalla
CS 5565—Lab 4

(e)



(f)

